

A Soft Introduction to Covariance Matrix Adaptation Evolution Strategy (CMA-ES)

From Gradient Descents to CMA-ES

Optimization Algorithm Team

Seangleng Khe (Cam)

Aphinan Photun (Tha)

Phoebe Jean Domen (Fil)

Luong Minh Tuan (Vie)

Bui Huynh Xuan Quyen (Vie)

Instructor

Prof. Gaël Raoul

September 11, 2026

SEAMS School 2026: Applying Mathematics to Real-World Applications
HO Chi Minh City University of Science, Vietnam

Introduction

The objective of the study is to understand the progression

gradient descent → random search → evolution strategy → CMA-ES

After: we shall see how important it is to adapt a full covariance matrix can make the algorithm nearly insensitive to rotations and anisotropic scalings of the search landscape.

Modified Objective: **Visit CMA-ES**


Optimization with Gradient Descents

Suppose we want to solve

$$\min_{x \in D} f(x)$$

If f is **smooth**, we use
gradient information

$$x_{n+1} = x_n - \sigma \nabla f(x_n)$$



Search information

GD uses **gradient information** to search for the optimizer.




Figure 1: 5D Rugged Landscape ¹

¹Hansen [2019]

Random Search

How can we reach the optimizer?

Random Search ²

With a starting point $m \Rightarrow$ generate random points around $m \Rightarrow$ calculate the function values f .

Search information

The **random search** uses a point m to generate **random points** based on probability distribution and calculate **the fitness value (function value)**.

Drawback: No directional learning/adaptation.

This is as stated " Finding a needle in an n-dimensional haystack " ³.

²White [1970]

³Korn [1966]

Evolution Strategy: (μ, λ) -ES

A simulation is here!

Search Information

(μ, λ) -ES uses a mean m generate to λ random points (called *offsprings*) and selects μ points (called *parents*), recombine to generate a new mean, and repeat the process.




Figure 2: Comparison between random search with uniform distribution and (μ, λ) -ES.

III-Conditioned Problems

(μ, λ) -ES is not enough...

- One observation is that the algorithm is inefficient when the successful search directions from a thin elongated ellipsoid.





Figure 3: Function f has anisotropic level set. (Steep valey in x_2 and flat in x_1).

Preferable search trajectory



Drawback:

Have a direction, but ineffective with ill-conditioned problems.

Covariance Matrix Adaptation Evolution Strategy (CMA-ES)

Search Information

CMA-ES = (μ, λ) -ES + learning/adaptation

(μ, λ) -ES generates points

$$X \sim \mathcal{N}(m, \sigma^2 I).$$

CMA-ES generates points

$$X \sim \mathcal{N}(m, \sigma^2 C)$$

where C is symmetric positive definite.

(μ, λ) -ES generates points

$$X \sim \mathcal{N}(m, \sigma^2 I).$$

CMA-ES generates points

$$X \sim \mathcal{N}(m, \sigma^2 C)$$

where C is symmetric positive definite.

What does the covariance matrix C do here?

A simulation is here!

- Directions in which successful steps have repeatedly occurred should receive larger variance in the future generations.




Figure 4: Function f has anisotropic level set. (Steep valey in x_2 and flat in x_1).

Search Information

CMA-ES learns not only where to search though the starting point m but also how to search through C .

What information that C gives?⁴

$$C_{n+1} = (1 - c_\mu) C_n + c_\mu \frac{1}{\sigma(i)^2} C_{n+1}^\mu$$

⁴Aguirre et al. [2026]

What information that C gives?⁴

$$C_{n+1} = \left(\underbrace{1 - c_\mu}_{\text{learning rate}} \right) C_n + c_\mu \frac{1}{\sigma(i)^2} \underbrace{C_{n+1}^\mu}_{\text{reproduce the best candidates}}$$

⁴Aguirre et al. [2026]

What information that C gives?⁴

$$C_{n+1} = \underbrace{\left(\underbrace{1 - c_\mu}_{\text{learning rate}} \right) C_n + c_\mu \frac{1}{\sigma(i)^2} \underbrace{C_{n+1}^\mu}_{\text{reproduce the best candidates}}}_{\text{memory process}}$$

⁴Aguirre et al. [2026]

There's something more...

We can control the step size σ ⁴

$$\ln \sigma_{n+1} = \ln \sigma_n + \frac{c_\sigma}{d_\sigma} \left(\frac{\|p_{n+1}^\sigma\|}{E \|\mathcal{N}(0, I)\|} - 1 \right)$$

⁴Aguirre et al. [2026]

There's something more...

We can control the step size σ^4

$$\ln \sigma_{n+1} = \ln \sigma_n + \frac{c_\sigma}{d_\sigma} \left(\underbrace{\frac{\|p_{n+1}^\sigma\|}{E \|\mathcal{N}(0, I)\|}}_{!!!} - 1 \right)$$

⁴Aguirre et al. [2026]

There's something more...

We can control the step size σ^4

$$\underbrace{\ln \sigma_{n+1} = \ln \sigma_n + \frac{c_\sigma}{d_\sigma} \left(\underbrace{\frac{\|p_{n+1}^\sigma\|}{E \|\mathcal{N}(0, I)\|} - 1}_{!!!} \right)}_{\text{memory process}}$$

⁴Aguirre et al. [2026]


The End of CMA-ES!!!

The End of CMA-ES!!!

... from me

Gradient-based Opt.

Derivative-Free Opt.



Gradient Methods

gradient/derivative

Random Search

generate random points

(μ, λ) -ES

random search +
repetition

CMA-ES

learning/adaptation (non-
trivial)

Presented Methods



CMA-ES vs Gradient Descent vs Quasi-Newton

Five landscapes with known minimizer $\mathbf{x}^*$, tested in $d = 2$ and $d = 10$:

1. **Sphere** $|\mathbf{x}|^2$: isotropic, convex
2. **Ellipsoid** $\sum_i 10^{6i/(d-1)} x_i^2$: separable, condition number 10^6 .
3. **Rotated ellipsoid** $\mathbf{x}^T \mathbf{Q}^T \mathbf{\Lambda} \mathbf{Q} \mathbf{x}$: same spectrum, principal axes no longer aligned with the coordinates.
4. **Rosenbrock**: non-convex but *smooth* and unimodal; a narrow valley whose curvature changes along its length.
5. **Rastrigin**: **multimodal**, **local minima** on a unit grid; tests global behaviour rather than local convergence.

Isotropic Quadratic Problem

Sphere $|x|^2$ ($d = 2$)



Convergence Analysis: $\|\mathbf{x}_n - \mathbf{x}^*\|$






Figure 5: Convergence Analysis on Isotropic Quadratic Problem

III-Conditioned Quadratic Problem

Separable ellipsoid ($d = 2$)



Convergence Analysis: $\|\mathbf{x}_n - \mathbf{x}^*\|$





Figure 6: Convergence Analysis on Ill-Conditioned Quadratic Problem

Rotated Ellipsoid Problem



Convergence Analysis: $\|\mathbf{x}_n - \mathbf{x}^*\|$






Figure 7: Convergence Analysis on Rotated Quadratic Problem

Rosenbrock Problem

Rosenbrock (curved valley, non-convex) ($d = 2$)



Convergence Analysis: $\|\mathbf{x}_n - \mathbf{x}^*\|$






Figure 8: Convergence Analysis on Rosenbrock Problem

Rastrigin Problem

Rastrigin (multimodal) ($d = 2$)



Convergence Analysis: $\|\mathbf{x}_n - \mathbf{x}^*\|$




Figure 9: Convergence Analysis on Rastrigin Problem

Success Rate: Runs Reaching $\|\mathbf{x}_n - \mathbf{x}^*\| < 10^{-6}$

Table: Percentage of the 50 runs reaching the target accuracy 10^{-6} within the budget of 2×10^4 function evaluations. Ill-conditioning and rotation break GD at $d = 10$; only Rastrigin resists all three methods.

Function	$d = 2$			$d = 10$		
	GD	BFGS	CMA-ES	GD	BFGS	CMA-ES
Sphere	100%	100%	100%	100%	100%	100%
Ellipsoid	100%	100%	100%	0%	100%	100%
Rotated ellipsoid	100%	100%	100%	0%	100%	100%
Rosenbrock	100%	100%	100%	0%	84%	88%
Rastrigin	2%	12%	24%	0%	0%	0%

At $d = 2$ every method solves every landscape except Rastrigin. Dimension is what separates them: GD fails on all four hard landscapes at $d = 10$, while CMA-ES matches BFGS everywhere and is slightly ahead on Rosenbrock.

Overall Evaluation

When Gradient Descent is Better

- **Isotropic & well-conditioned problems** (e.g., Sphere).
- **Execution Speed:** Extremely fast due to exact analytical gradient steps.
- Highly scalable to large dimensions **if** the landscape is smooth and convex.

When CMA-ES is Better

- **Ill-conditioned & correlated landscapes** (e.g., Rotated Ellipsoid): Effectively adapts its covariance matrix to match the problem's topography.
- **Non-convex & multimodal problems** (e.g., Rosenbrock, Rastrigin): Its stochastic nature helps escape narrow valleys and shallow local minima.
- Does not require gradient information (derivative-free optimization).

References

- Nikolaus Hansen. How to evolve gradient descent into evolution strategies and CMA-ES. Presentation slides, Symposium on Evolutionary Algorithms: Back to the Future and Beyond, Delft, Netherlands, September 2019. URL <http://www.cmap.polytechnique.fr/~nikolaus.hansen/cmaes-from-gradient-descent-delft2019.pdf>.
- R.C. White. *A survey of random methods for parameter optimization*. EUT report. E, Fac. of Electrical Engineering. Technische Hogeschool Eindhoven, 1970.
- Granino Arthur Korn. *Random-Process Simulation and Measurements*. McGraw-Hill, New York, 1966.
- Nahum Aguirre, Erik Cuevas, Alberto Luque-Chang, Oscar Barba-Toscano, and Mario Vasquez-Franco. Covariance matrix adaptation evolution strategy (CMA-ES): A comprehensive survey of variants and hybridizations. *Archives of Computational Methods in Engineering*, 2026. doi: 10.1007/s11831-026-10557-z.

**Thank you for your attentions!!!
:-)**